




Held-out predictor parity
 $R^2=0.775$ MAE=0.094 90% coverage=0.94



Pareto comparison: CVAE vs random vs GA



Encoder neighbours share elements; mean=0.92



Active-learning recovery test

Hold out the top 20 STY rows.

Retrain a fresh ActivityHead on the rest.

Recovered in top-50: 8 / 20 (40%).

Top-decile latent coherence

Top decile rows have on average 48% of their 5 nearest latent also in the top decile.